


## Homework Assignment for Chapter 12

### 12.1 The Hepatic Injury

#### a) Training and Testing set.

Given the classification imbalance, stratified random sampling method would be the best approach in splitting the data into training and testing set.

#### b) Classification statistic



**Fig 1:** Distribution of the different classes in the response variable of the dataset

Given the distribution has more than 2 classes, kappa and accuracy are the best classification statistics. I would choose Kappa as a classification statistic to optimize for this exercise.

## c) Training and Testing

### Logistic Regression

```
225 samples
96 predictor
3 classes: 'Mild', 'None', 'Severe'

No pre-processing
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 169, 169, 169, 169, 169, 169, ...
Resampling results across tuning parameters:
```

decay	logLoss	AUC	prAUC	Accuracy	Kappa	Mean_F1	Mean_Sensitivity	Mean_Specificity	Mean_Pos_Pred_Value
0e+00	19.849982	0.5436728	0.2386227	0.3921429	0.037535713	0.3522554	0.3660427	0.6790519	0.3595268
1e-04	11.352295	0.5473951	0.3362843	0.4000000	0.038452093	0.3523776	0.3673892	0.6784508	0.3571626
1e-01	2.071993	0.5385078	0.3499925	0.4278571	0.003155277	0.3747908	0.3541434	0.6652134	0.3638287
Mean_Neg_Pred_Value Mean_Precision Mean_Recall Mean_Detection_Rate Mean_Balanced_Accuracy									
0.6783829		0.3595268	0.3660427	0.1307143		0.5225473			
0.6789135		0.3571626	0.3673892	0.1333333		0.5229200			
0.6639770		0.3638287	0.3541434	0.1426190		0.5096784			

Kappa was used to select the optimal model using the largest value.  
The final value used for the model was decay = 1e-04.

### Confusion Matrix and Statistics

	Reference		
Prediction	Mild	None	Severe
Mild	12	11	5
None	12	5	0
Severe	5	5	1

### Overall Statistics

Accuracy : 0.3214  
95% CI : (0.2029, 0.4596)  
No Information Rate : 0.5179  
P-Value [Acc > NIR] : 0.9990


Kappa : -0.1194

McNemar's Test P-Value : 0.1686

### Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.4138	0.23810	0.16667
Specificity	0.4074	0.65714	0.80000
Pos Pred Value	0.4286	0.29412	0.09091
Neg Pred Value	0.3929	0.58974	0.88889
Prevalence	0.5179	0.37500	0.10714
Detection Rate	0.2143	0.08929	0.01786
Detection Prevalence	0.5000	0.30357	0.19643
Balanced Accuracy	0.4106	0.44762	0.48333

> |



**Fig 2:** Plot of Kappa vs Weigh Decay for Logistic Regression

## LDA

### Linear Discriminant Analysis

225 samples  
96 predictor  
3 classes: 'Mild', 'None', 'Severe'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 169, 169, 169, 169, 169, 169, ...

Resampling results:

logLoss	AUC	prAUC	Accuracy	Kappa	Mean_F1	Mean_Sensitivity	Mean_Specificity	Mean_Pos_Pred_Value
4.646426	0.5413375	0.3591677	0.4214286	0.04851528	0.3594966	0.3681554	0.6838434	0.3624741
Mean_Neg_Pred_Value	Mean_Precision	Mean_Recall	Mean_Detection_Rate	Mean_Balanced_Accuracy				
0.6833069	0.3624741	0.3681554	0.1404762	0.5259994				

### Confusion Matrix and Statistics

	Reference		
Prediction	Mild	None	Severe
Mild	18	7	2
None	10	12	1
Severe	1	2	3

### Overall Statistics

Accuracy : 0.5893  
95% CI : (0.4498, 0.719)  
No Information Rate : 0.5179  
P-Value [Acc > NIR] : 0.1747

Kappa : 0.2977

Mcnemar's Test P-Value : 0.7539

### Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.6207	0.5714	0.50000
Specificity	0.6667	0.6857	0.94000
Pos Pred Value	0.6667	0.5217	0.50000
Neg Pred Value	0.6207	0.7273	0.94000
Prevalence	0.5179	0.3750	0.10714
Detection Rate	0.3214	0.2143	0.05357
Detection Prevalence	0.4821	0.4107	0.10714
Balanced Accuracy	0.6437	0.6286	0.72000

# PLSDA

Partial Least Squares

225 samples  
96 predictor  
3 classes: 'Mild', 'None', 'Severe'

Pre-processing: centered (96), scaled (96)  
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)  
Summary of sample sizes: 169, 169, 169, 169, 169, 169, ...  
Resampling results across tuning parameters:

ncomp	logLoss	AUC	prAUC	Accuracy	Kappa	Mean_F1	Mean_Sensitivity	Mean_Specificity	Mean_Pos_Pred_Value
1	1.029974	0.5314927	0.3520355	0.5100000	0.052719936	NaN	0.3548987	0.6827937	NaN
2	1.051596	0.5233578	0.3448435	0.4864286	0.019030137	NaN	0.3414778	0.6723668	NaN
3	1.064066	0.5300979	0.3462342	0.4878571	0.035304728	NaN	0.3469513	0.6785467	NaN
4	1.070227	0.5384948	0.3513280	0.4750000	0.019669134	0.3377176	0.3411385	0.6735041	0.3459017
5	1.078944	0.5461276	0.3515789	0.4700000	0.012874258	0.4297998	0.3407444	0.6710688	0.3967828
6	1.088574	0.5437108	0.3514548	0.4864286	0.046249714	0.4160606	0.3601204	0.6816353	0.4654380
7	1.108331	0.5351090	0.3452783	0.4628571	0.006963565	0.3772711	0.3395074	0.6693079	0.3774511
8	1.120990	0.5352470	0.3453930	0.4678571	0.024051798	0.4025025	0.3460536	0.6750293	0.3587227
9	1.126809	0.5363608	0.3494551	0.4678571	0.030788391	0.3868731	0.3497537	0.6776395	0.3586208
10	1.137619	0.5391484	0.3520954	0.4735714	0.043787416	0.3886049	0.3536070	0.6821630	0.3645169
11	1.146572	0.5381211	0.3494566	0.4671429	0.038185786	0.3908116	0.3510564	0.6803302	0.3571913
12	1.153089	0.5402534	0.3524448	0.4685714	0.045767663	0.3879850	0.3530268	0.6829108	0.3454018
13	1.156581	0.5411761	0.3536293	0.4707143	0.051486885	0.3793724	0.3547564	0.6850765	0.3484363
14	1.157705	0.5453142	0.3588207	0.4792857	0.070475368	0.3979530	0.3651998	0.6914610	0.3532554
15	1.161970	0.5464935	0.3587136	0.4757143	0.069357772	0.3935199	0.3643021	0.6908063	0.3515199
Mean_Neg_Pred_Value Mean_Precision Mean_Recall Mean_Detection_Rate Mean_Balanced_Accuracy									
0.6900099		NaN	0.3548987	0.1700000	0.5188462				
0.6744536		NaN	0.3414778	0.1621429	0.5069223				
0.6829147		NaN	0.3469513	0.1626190	0.5127490				
0.6758845	0.3459017		0.3411385	0.1583333	0.5073213				
0.6729549	0.3967828		0.3407444	0.1566667	0.5059066				
0.6851454	0.4654380		0.3601204	0.1621429	0.5208778				
0.6713721	0.3774511		0.3395074	0.1542857	0.5044077				
0.6770840	0.3587227		0.3460536	0.1559524	0.5105415				
0.6795845	0.3586208		0.3497537	0.1559524	0.5136966				
0.6854422	0.3645169		0.3536070	0.1578571	0.5178850				
0.6827195	0.3571913		0.3510564	0.1557143	0.5156933				
0.6853472	0.3454018		0.3530268	0.1561905	0.5179688				
0.6875420	0.3484363		0.3547564	0.1569048	0.5199165				
0.6942946	0.3532554		0.3651998	0.1597619	0.5283304				
0.6929204	0.3515199		0.3643021	0.1585714	0.5275542				

Kappa was used to select the optimal model using the largest value.  
The final value used for the model was ncomp = 14.

### Confusion Matrix and Statistics

Reference			
Prediction	Mild	None	Severe
Mild	17	11	3
None	12	9	2
Severe	0	1	1

### Overall Statistics


Accuracy : 0.4821  
95% CI : (0.3466, 0.6197)  
No Information Rate : 0.5179  
P-Value [Acc > NIR] : 0.7482

Kappa : 0.0677

McNemar's Test P-Value : 0.3371

### Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.5862	0.4286	0.16667
Specificity	0.4815	0.6000	0.98000
Pos Pred Value	0.5484	0.3913	0.50000
Neg Pred Value	0.5200	0.6364	0.90741
Prevalence	0.5179	0.3750	0.10714
Detection Rate	0.3036	0.1607	0.01786
Detection Prevalence	0.5536	0.4107	0.03571
Balanced Accuracy	0.5338	0.5143	0.57333



**Fig 3:** Plot of tuning parameter for PLSDA model

## Glmnet

```
glmnet

225 samples
96 predictor
3 classes: 'Mild', 'None', 'Severe'

Pre-processing: centered (96), scaled (96)
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 169, 169, 169, 169, 169, 169, ...
Resampling results across tuning parameters:
```

alpha	lambda	logLoss	AUC	prAUC	Accuracy	Kappa	Mean_F1
0.0	0.01000000	2.1117474	0.5503897	0.357011044	0.4614286	0.058379737	0.3930794
0.0	0.03111111	1.5789332	0.5490661	0.352513538	0.4778571	0.062190735	0.4072934
0.0	0.05222222	1.4012425	0.5493445	0.352489622	0.4778571	0.048091939	0.3916180
0.0	0.07333333	1.3065767	0.5498002	0.352194090	0.4792857	0.045051622	0.3899835
0.0	0.09444444	1.2460239	0.5497011	0.351767249	0.4821429	0.043447276	0.3920211
0.0	0.11555556	1.2034180	0.5501954	0.352261830	0.4871429	0.046545530	0.3916944
0.0	0.13666667	1.1715110	0.5507062	0.352427990	0.4892857	0.046390871	0.3893633
0.0	0.15777778	1.1467007	0.5509706	0.353066580	0.4935714	0.049037172	0.3674336
0.0	0.17888889	1.1266409	0.5507307	0.353860542	0.4892857	0.038228036	0.3663536
0.0	0.20000000	1.1101390	0.5506610	0.355263963	0.4885714	0.032932254	0.3750842
0.1	0.01000000	2.0180586	0.5510356	0.355882434	0.4671429	0.062880758	0.4029170
0.1	0.03111111	1.4439662	0.5471923	0.351539680	0.4721429	0.042018618	0.3886079
0.1	0.05222222	1.2662955	0.5472742	0.352780964	0.4785714	0.038513542	0.3853954
0.1	0.07333333	1.1757593	0.5493696	0.355813675	0.4842857	0.038512670	0.4074613
0.1	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
0.1	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
0.1	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
0.1	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
0.1	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
0.1	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN
0.2	0.01000000	1.9175937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
0.2	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
0.2	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
0.2	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
0.3	0.01000000	1.8325937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
0.3	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
0.3	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
0.3	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
0.3	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
0.3	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
0.3	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
0.3	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
0.3	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
0.3	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN
0.4	0.01000000	1.8325937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
0.4	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
0.4	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
0.4	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
0.4	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
0.4	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
0.4	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
0.4	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
0.4	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
0.4	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN
0.5	0.01000000	1.8325937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
0.5	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
0.5	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
0.5	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
0.5	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
0.5	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
0.5	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
0.5	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
0.5	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
0.5	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN
0.6	0.01000000	1.8325937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
0.6	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
0.6	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
0.6	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
0.6	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
0.6	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
0.6	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
0.6	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
0.6	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
0.6	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN
0.7	0.01000000	1.8325937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
0.7	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
0.7	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
0.7	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
0.7	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
0.7	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
0.7	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
0.7	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
0.7	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
0.7	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN
0.8	0.01000000	1.8325937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
0.8	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
0.8	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
0.8	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
0.8	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
0.8	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
0.8	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
0.8	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
0.8	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
0.8	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN
0.9	0.01000000	1.8325937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
0.9	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
0.9	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
0.9	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
0.9	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
0.9	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
0.9	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
0.9	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
0.9	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
0.9	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN
1.0	0.01000000	1.8325937	0.5492235	0.355895550	0.4635714	0.053700776	0.4003238
1.0	0.03111111	1.3431075	0.5475500	0.353914575	0.4771429	0.042084080	0.3974343
1.0	0.05222222	1.1794419	0.5490482	0.355374039	0.4800000	0.029021478	0.4081553
1.0	0.07333333	1.1402571	0.5473611	0.355763000	0.4878571	0.027515700	...
1.0	0.09444444	1.1213088	0.5503700	0.356646898	0.4828571	0.024652295	0.3946197
1.0	0.11555556	1.0855167	0.5479795	0.355169091	0.4864286	0.022796884	NaN
1.0	0.13666667	1.0601313	0.5464738	0.354263181	0.4878571	0.019598612	NaN
1.0	0.15777778	1.0407978	0.5459843	0.354547163	0.4942857	0.025373068	NaN
1.0	0.17888889	1.0258990	0.5446728	0.353850814	0.4928571	0.017601509	NaN
1.0	0.20000000	1.0139374	0.5430350	0.352752720	0.5000000	0.028572519	NaN

```
[ reached getOption("max.print") -- omitted 8 rows ]

Kappa was used to select the optimal model using the largest value.
The final values used for the model were alpha = 0 and lambda = 0.
```

### Confusion Matrix and Statistics

	Reference		
Prediction	Mild	None	Severe
Mild	17	9	3
None	12	11	2
Severe	0	1	1

### Overall Statistics

Accuracy : 0.5179  
 95% CI : (0.3803, 0.6534)  
 No Information Rate : 0.5179  
 P-Value [Acc > NIR] : 0.5537

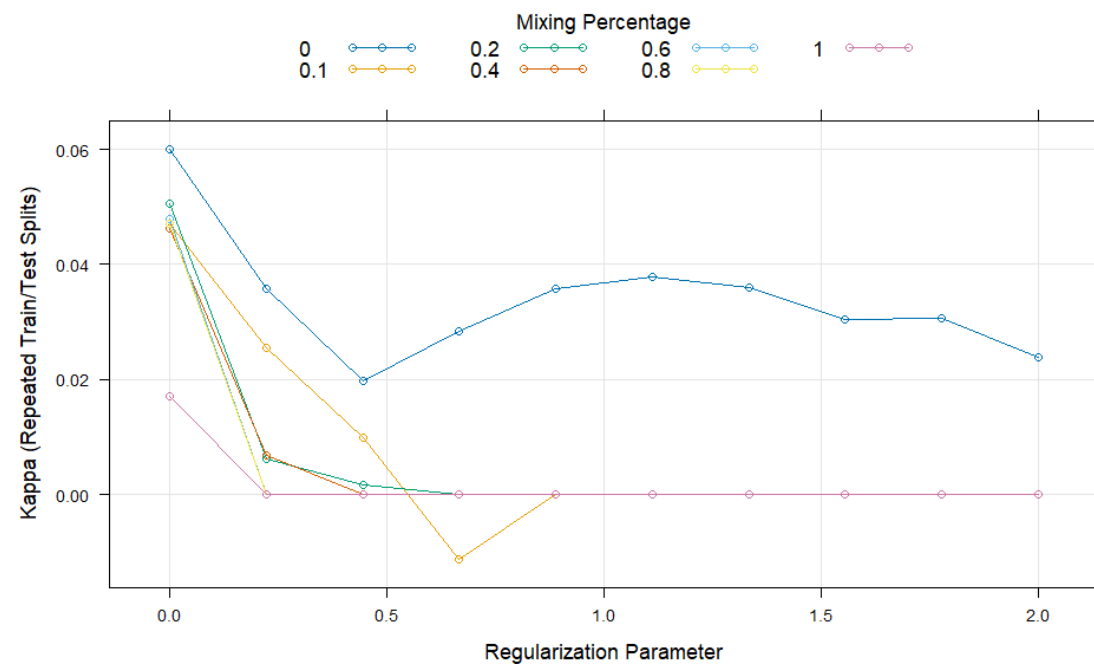
Kappa : 0.1399

McNemar's Test P-Value : 0.2883

### Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.5862	0.5238	0.16667
Specificity	0.5556	0.6000	0.98000
Pos Pred Value	0.5862	0.4400	0.50000
Neg Pred Value	0.5556	0.6774	0.90741
Prevalence	0.5179	0.3750	0.10714
Detection Rate	0.3036	0.1964	0.01786
Detection Prevalence	0.5179	0.4464	0.03571
Balanced Accuracy	0.5709	0.5619	0.57333

### Penalized Logistic Tuning Parameters




**Fig 4:** Plot of tuning parameter for glmnet model



Model	Best Tuning Parameter	Training		Testing	
		Kappa	Accuracy	Kappa	Accuracy
<b>Logistic Regression</b>	Decay=1e-04	0.03845	0.4000	-0.1194	0.3214
<b>Linear Discriminant Analysis</b>	N/A	0.04851528	0.4214286	0.2977	0.5893
<b>PLSDA</b>	ncomp=14	0.0704754	0.4792857	0.0677	0.4821
<b>glmnet</b>	alpha=0, lambda=0.0	0.06000	0.4621429	0.1399	0.5179

Looking at the results from the table, the models exhibit different performances. On the training, PLSDA has the best performance with an accuracy of 47.93% and kappa of 0.07048. On the testing, LDA has the best performance with an accuracy of 58.93% and kappa of 0.2977. However, LDA performs poorly on the training while PLSDA performs poorly on testing. The best model for predicting injury based on biological factors is glmnet as it performs better on both training and testing with Accuracy of 46.21% and kappa of 0.06 on training and accuracy of 51.79% and kappa of 0.1399 on testing. Glmnet was the second-best performing model on both training and testing while LDA and PLSDA fluctuated in performance.

#### d) Five important predictors



**Fig 5:** Important features for glmnet model

## 12.3 MLC++ software package

### a) Explore the data

The 'mlc\_churn' dataset has 5000 observations and 20 predictors

#### NearZeroVariance

```
> nZv <- nearZeroVar(mlc_churn)
> length(nZv)
[1] 1
> |
```

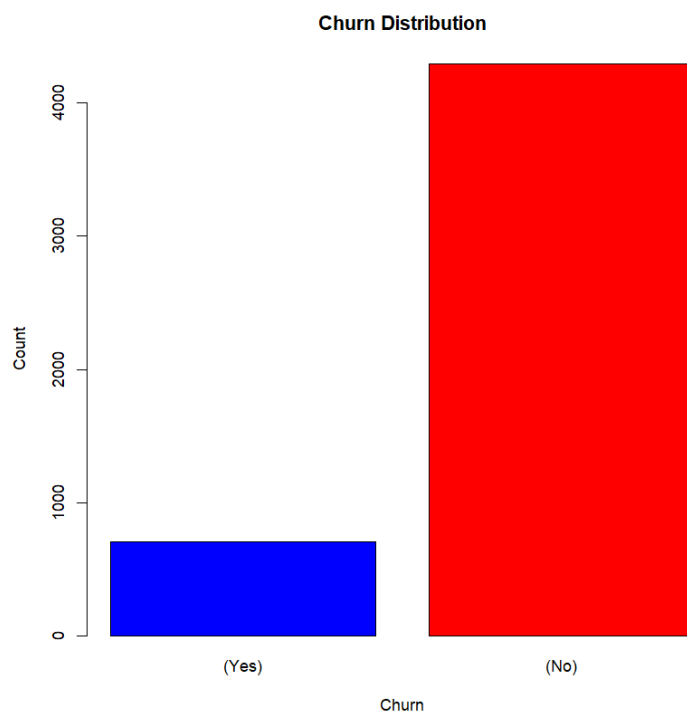
From the dataset, one predictor was found to have near zero variance. The predictor was removed.

#### Distribution for response variable

```
> counts

yes  no
707 4293
> percentage <- prop.table(counts) * 100
> percentage


yes  no
14.14 85.86
```



**Fig 6:** Churn distribution

The response variable was churn, with two classes('Yes', 'No'). 14.14% of the variables were under 'Yes' class and 85.86% were under the 'No' class.










### Barplot for Categorical pred

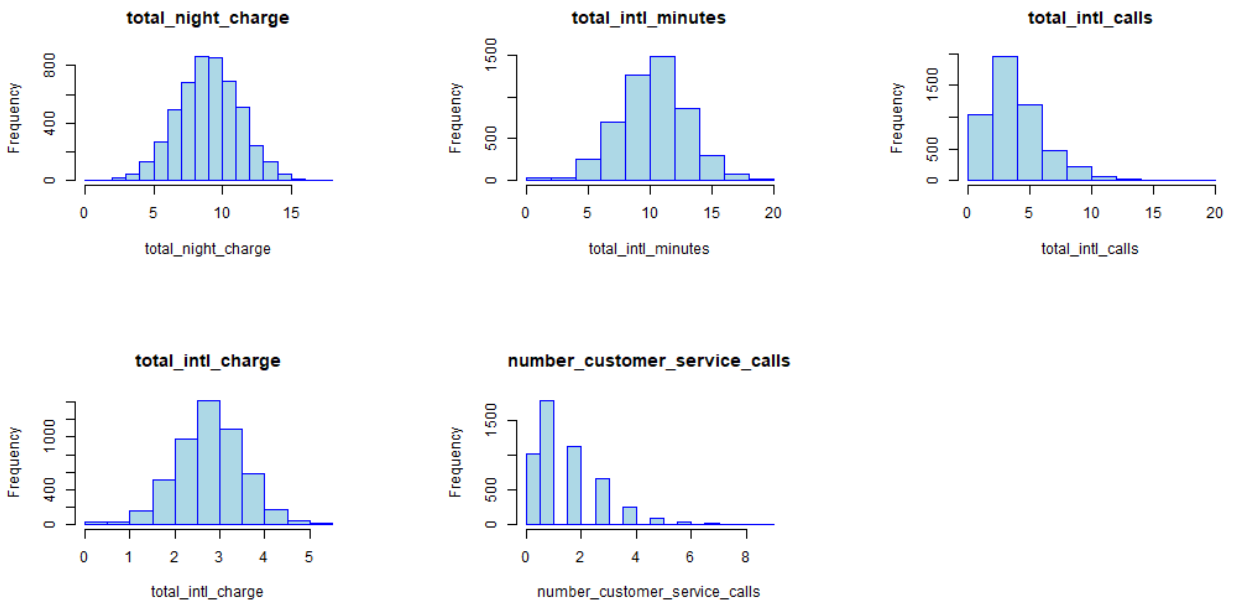


**Fig 7:** Distribution for categorical predictors

Looking at the categorical predictor, State was a degenerate predictor. It shows little to no variability across different observations. This predictor was removed before modelling.

## Histogram for Numerical predictors






**Fig 8:** Histogram plots for numerical predictors


The histogram plots show the distribution of the numerical predictors. From the plots, some numerical predictors had skewness and the need to remove the skewness was evident.

## Boxplot distribution for Numerical predictors to check on outliers


account\_length Distribution




total\_day\_minutes Distribution




total\_day\_calls Distribution




total\_day\_charge Distribution




total\_eve\_minutes Distribution




total\_eve\_calls Distribution




total\_eve\_charge Distribution

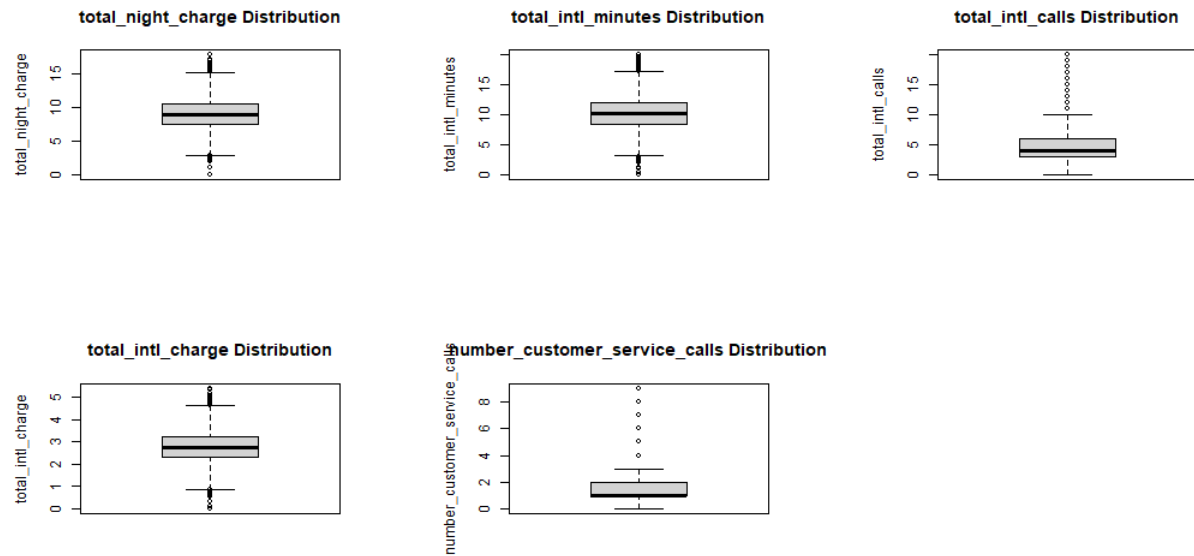


total\_night\_minutes Distribution



total\_night\_calls Distribution






**Fig 9:** Boxplot distribution for numerical predictors

From the boxplots, some numerical predictors had outliers.

## Correlation

```
> highCorr
[1] 8 13 7 4
> length(highCorr)
[1] 4
```

## Correlation plot for numerical predictors




**Fig 10:** Correlation plots for numerical predictors

Four predictors were highly correlated and the need to remove them was evident.



## Correlation plot after removing highly correlated predictors



**Fig 11:** Correlation plot after removing highly correlated predictors

## Preprocessing and transformation

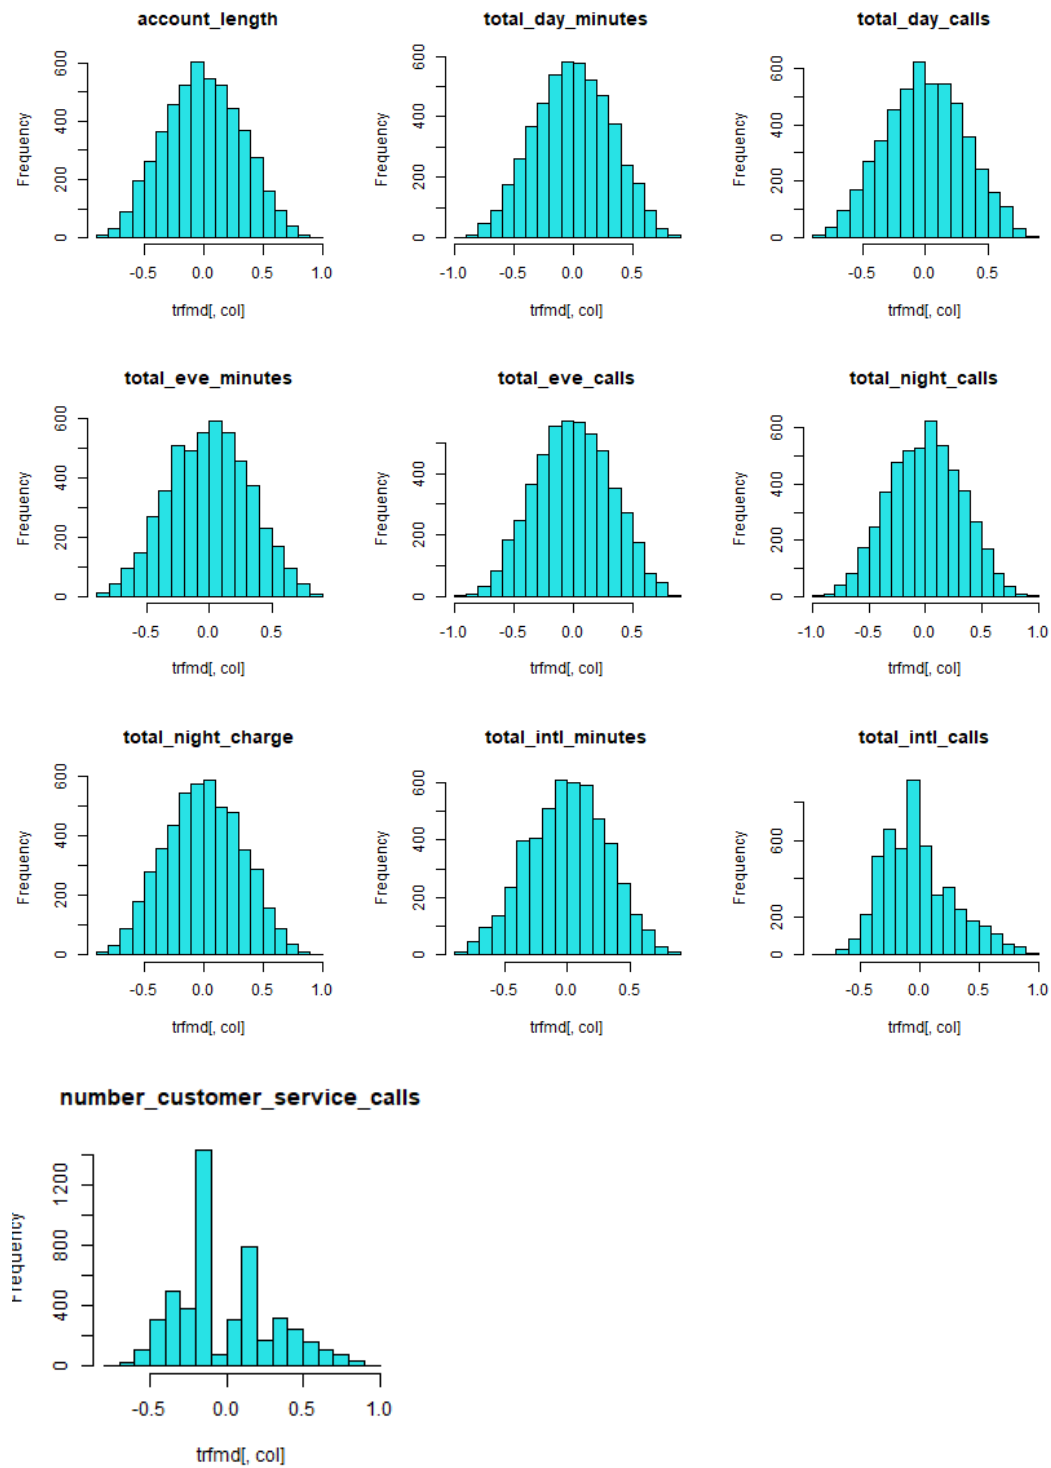
Created from 5000 samples and 10 variables

Pre-processing:

- Box-Cox transformation (1)
- centered (10)
- ignored (0)
- scaled (10)
- spatial sign transformation (10)

Lambda estimates for Box-Cox transformation:  
0.9

## Numerical predictors after transformation



**Fig 12:** Distribution of numerical predictors after transformation

## b) Evaluation of model

Area under the curve(AUC-ROC) was used to evaluate the effectiveness of the models

## c) Training and testing models

### Logistic regression

```
Generalized Linear Model
```

```
4001 samples
 14 predictor
 2 classes: 'yes', 'no'
```

```
No pre-processing
```

```
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
```

```
Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...
```

```
Resampling results:
```

ROC	Sens	Spec
0.8085635	0.177305	0.9765035

```
Confusion Matrix and Statistics
```

	Reference	
Prediction	yes	no
yes	25	19
no	116	839

```
Accuracy : 0.8649
```

```
95% CI : (0.8421, 0.8855)
```

```
No Information Rate : 0.8589
```

```
P-Value [Acc > NIR] : 0.3115
```

```
Kappa : 0.2178
```

```
Mcnemar's Test P-Value : <2e-16
```

```
Sensitivity : 0.17730
```

```
Specificity : 0.97786
```

```
Pos Pred Value : 0.56818
```

```
Neg Pred Value : 0.87853
```


```
Prevalence : 0.14114
```

```
Detection Rate : 0.02503
```

```
Detection Prevalence : 0.04404
```

```
Balanced Accuracy : 0.57758
```

```
'Positive' Class : yes
```



**Fig 13:** AUC-ROC curve for Logistic regression after prediction

## Linear Discriminant Analysis

Linear Discriminant Analysis

4001 samples  
 14 predictor  
 2 classes: 'yes', 'no'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...

Resampling results:

ROC	Sens	Spec
0.8110926	0.218156	0.9627506

## Confusion Matrix and Statistics

	Reference	
Prediction	yes	no
yes	34	25
no	107	833


Accuracy : 0.8679  
95% CI : (0.8453, 0.8883)  
No Information Rate : 0.8589  
P-Value [Acc > NIR] : 0.2212

Kappa : 0.28

Mcnemar's Test P-Value : 1.787e-12

Sensitivity : 0.24113  
Specificity : 0.97086  
Pos Pred Value : 0.57627  
Neg Pred Value : 0.88617  
Prevalence : 0.14114  
Detection Rate : 0.03403  
Detection Prevalence : 0.05906  
Balanced Accuracy : 0.60600

'Positive' Class : yes



**Fig 14:** AUC-ROC curve for LDA after prediction

## PLSDA

Partial Least Squares

4001 samples

14 predictor

2 classes: 'yes', 'no'

No pre-processing

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...

Resampling results across tuning parameters:

ncomp	ROC	Sens	Spec
1	0.8014561	0.04283688	0.9950583
2	0.8103629	0.06893617	0.9913287
3	0.8109812	0.07347518	0.9902564
4	0.8110840	0.08028369	0.9895105
5	0.8110909	0.08028369	0.9895105
6	0.8110916	0.08028369	0.9895105
7	0.8110926	0.08028369	0.9895105
8	0.8110926	0.08028369	0.9895105
9	0.8110926	0.08028369	0.9895105
10	0.8110926	0.08028369	0.9895105
11	0.8110926	0.08028369	0.9895105
12	0.8110926	0.08028369	0.9895105
13	0.8110926	0.08028369	0.9895105
14	0.8110926	0.08028369	0.9895105

ROC was used to select the optimal model using the largest value.

The final value used for the model was ncomp = 7.

## Confusion Matrix and Statistics

	Reference	
Prediction	yes	no
yes	15	9
no	126	849


Accuracy : 0.8649  
95% CI : (0.8421, 0.8855)  
No Information Rate : 0.8589  
P-Value [Acc > NIR] : 0.3115

Kappa : 0.1468

Mcnemar's Test P-Value : <2e-16

Sensitivity : 0.10638  
Specificity : 0.98951  
Pos Pred Value : 0.62500  
Neg Pred Value : 0.87077  
Prevalence : 0.14114  
Detection Rate : 0.01502  
Detection Prevalence : 0.02402  
Balanced Accuracy : 0.54795

'Positive' Class : yes



**Fig 15:** AUC-ROC curve for PLSDA after prediction

## GLMNET

glmnet

4001 samples

14 predictor

2 classes: 'yes', 'no'

Pre-processing: centered (14), scaled (14)

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...

Resampling results across tuning parameters:

alpha	lambda	ROC	Sens	Spec
0.0	0.01000000	0.8094001	0.1407092199	0.9829371
0.0	0.03111111	0.8102475	0.0856737589	0.9893240
0.0	0.05222222	0.8106373	0.0581560284	0.9926807
0.0	0.07333333	0.8108304	0.0419858156	0.9960373
0.0	0.09444444	0.8109871	0.0255319149	0.9975758
0.0	0.11555556	0.8109828	0.0130496454	0.9987879
0.0	0.13666667	0.8110060	0.0073758865	0.9994872
0.0	0.15777778	0.8109785	0.0011347518	0.9996270
0.0	0.17888889	0.8109451	0.0000000000	0.9998135
0.0	0.20000000	0.8109084	0.0000000000	1.0000000
0.1	0.01000000	0.8097661	0.1370212766	0.9835431
0.1	0.03111111	0.8109653	0.0765957447	0.9901166
0.1	0.05222222	0.8113016	0.0462411348	0.9941725
0.1	0.07333333	0.8111210	0.0241134752	0.9976690
0.1	0.09444444	0.8105890	0.0070921986	0.9993473
0.1	0.11555556	0.8094708	0.0014184397	0.9997203
0.1	0.13666667	0.8080656	0.0000000000	1.0000000
0.1	0.15777778	0.8064894	0.0000000000	1.0000000
0.1	0.17888889	0.8049189	0.0000000000	1.0000000
0.1	0.20000000	0.8035358	0.0000000000	1.0000000
0.2	0.01000000	0.8100094	0.1336170213	0.9842890
0.2	0.03111111	0.8109795	0.0700709220	0.9909557
0.2	0.05222222	0.8102052	0.0346099291	0.9953380
0.2	0.07333333	0.8075954	0.0096453901	0.9991608
0.2	0.09444444	0.8044461	0.0014184397	0.9998135
0.2	0.11555556	0.8012852	0.0000000000	1.0000000
0.2	0.13666667	0.7975043	0.0000000000	1.0000000
0.2	0.15777778	0.7945336	0.0000000000	1.0000000
0.2	0.17888889	0.7911762	0.0000000000	1.0000000
0.2	0.20000000	0.7870066	0.0000000000	1.0000000
0.4	0.01000000	0.8102729	0.1248226950	0.9850350
0.4	0.03111111	0.8090142	0.0519148936	0.9924476
0.4	0.05222222	0.8031120	0.0110638298	0.9984615



0.4	0.03111111	0.8090142	0.0519148936	0.9924476
0.4	0.05222222	0.8031120	0.0110638298	0.9984615
0.4	0.07333333	0.7961259	0.0002836879	1.0000000
0.4	0.09444444	0.7892751	0.0000000000	1.0000000
0.4	0.11555556	0.7836820	0.0000000000	1.0000000
0.4	0.13666667	0.7767690	0.0000000000	1.0000000
0.4	0.15777778	0.7461396	0.0000000000	1.0000000
0.4	0.17888889	0.6495169	0.0000000000	1.0000000
0.4	0.20000000	0.5909002	0.0000000000	1.0000000
0.6	0.01000000	0.8102263	0.1129078014	0.9860606
0.6	0.03111111	0.8044884	0.0365957447	0.9943124
0.6	0.05222222	0.7947462	0.0011347518	0.9998135
0.6	0.07333333	0.7845906	0.0000000000	1.0000000
0.6	0.09444444	0.7727855	0.0000000000	1.0000000
0.6	0.11555556	0.6681506	0.0000000000	1.0000000
0.6	0.13666667	0.5817125	0.0000000000	1.0000000
0.6	0.15777778	0.5000000	0.0000000000	1.0000000
0.6	0.17888889	0.5000000	0.0000000000	1.0000000
0.6	0.20000000	0.5000000	0.0000000000	1.0000000
0.8	0.01000000	0.8100967	0.1043971631	0.9869464
0.8	0.03111111	0.7996961	0.0224113475	0.9962704
0.8	0.05222222	0.7857829	0.0000000000	1.0000000
0.8	0.07333333	0.7671200	0.0000000000	1.0000000
0.8	0.09444444	0.6135949	0.0000000000	1.0000000
0.8	0.11555556	0.5000000	0.0000000000	1.0000000
0.8	0.13666667	0.5000000	0.0000000000	1.0000000
0.8	0.15777778	0.5000000	0.0000000000	1.0000000
0.8	0.17888889	0.5000000	0.0000000000	1.0000000
0.8	0.20000000	0.5000000	0.0000000000	1.0000000
1.0	0.01000000	0.8096652	0.0941843972	0.9875991
1.0	0.03111111	0.7947905	0.0099290780	0.9986946
1.0	0.05222222	0.7802476	0.0000000000	1.0000000
1.0	0.07333333	0.6333762	0.0000000000	1.0000000
1.0	0.09444444	0.5000000	0.0000000000	1.0000000
1.0	0.11555556	0.5000000	0.0000000000	1.0000000
1.0	0.13666667	0.5000000	0.0000000000	1.0000000
1.0	0.15777778	0.5000000	0.0000000000	1.0000000
1.0	0.17888889	0.5000000	0.0000000000	1.0000000
1.0	0.20000000	0.5000000	0.0000000000	1.0000000

ROC was used to select the optimal model using the largest value.  
The final values used for the model were alpha = 0.1 and lambda = 0.05222222.

## Confusion Matrix and Statistics

```

      Reference
Prediction yes  no
yes      10    3
no     131  855


      Accuracy : 0.8659
      95% CI   : (0.8432, 0.8864)
No Information Rate : 0.8589
P-Value [Acc > NIR] : 0.2798

      Kappa : 0.1086

McNemar's Test P-Value : <2e-16

      Sensitivity : 0.07092
      Specificity : 0.99650
      Pos Pred Value : 0.76923
      Neg Pred Value : 0.86714
      Prevalence : 0.14114
      Detection Rate : 0.01001
      Detection Prevalence : 0.01301
      Balanced Accuracy : 0.53371

      'Positive' Class : yes
```



**Fig 16:** AUC-ROC curve for PLSDA after prediction

<b>Model</b>	<b>Best Tuning Parameter</b>	<b>Training (AUC-ROC)</b>	<b>Testing (AUC-ROC)</b>
<b>Logistic Regression</b>	N/A	0.8085635	0.8211
<b>Linear Discriminant Analysis</b>	N/A	0.8110926	0.8287
<b>PLSDA</b>	ncomp=7	0.8110926	0.8287
<b>glmnet</b>	alpha=0.1, lambda=0.05222	0.8110316	0.8279

From the summary in the table above, PLSDA and LDA have similar results in terms of their performance on the training and testing sets, using area under the curve. The two models have a performance of 81.11% on training and 82.87% on testing. Glmnet is the closest model to the two with a performance of 81.103% on training and 82.79% on testing. The best model would be PLSDA as it is the joint performing model with LDA, but it has a tuning parameter with ncomp=7 as the best tuning parameter for PLSDA. This parameter can be adjusted to improve the performance of the model.

## Appendix

##Question 12.1

```
library(caret)
```

```
library(AppliedPredictiveModeling)
```

```
data(hepatic)
```

```
injury
```

```
table(injury)
```

```
barplot(table(injury), col=c('yellow','green','blue'),main='Class Distribution')
```

```
#part c: Nearzero & Corr
```

```
bio <- bio[, -nearZeroVar(bio)]
```

```
highCorrbio <- findCorrelation(cor(bio), cutoff=0.90)
```

```
bio<- bio[, -highCorrbio]
```

```
##split the data 80/20
```

```
install.packages("MLmetrics")
```

```
set.seed(100)
```

```
trainR <- createDataPartition(injury, p=0.8, list=FALSE)
```

```
X.train <- bio[trainR, ]
```

```
y.train <- injury[trainR]
```

```
X.test <- bio[-trainR, ]
```

```
y.test <- injury[-trainR]
```

```
ctrl<-trainControl(method='LGOCV',summaryFunction = multiClassSummary,classProbs =  
TRUE)
```

```
##Running models
#install.packages(c("glmnet", "pamr", "rms", "sparseLDA", "subselect"))
```

```
#####Logistic Regression#####
```

```
X.train <- as.data.frame(X.train)
```

```
set.seed(100)
```

```
bio_lr <- caret::train(X.train,
  y.train,
  method = 'multinom',
  metric = "Kappa",
  trControl = ctrl,
  )
```

```
bio_lr
```

```
plot(bio_lr)
```

```
pred_bio<-predict(bio_lr,X.test)
```

```
confusionMatrix(data=pred_bio,
  reference=y.test)
```

```
reducedRoc <- roc(response = lrReduced$pred$obs,
  predictor = lrReduced$pred$successful,
  levels = rev(levels(lrReduced$pred$obs)))
```

```
plot(reducedRoc, legacy.axes = TRUE)
```

```
auc(reducedRoc)
```

```
#####Linear Discriminant Analysis#####
```

```
library(MASS)
```

```
set.seed(100)
```

```
bio_lda <- caret::train(x = X.train,  
  y = y.train,  
  method = "lda",  
  metric = "Kappa",  
  trControl = ctrl)
```

```
bio_lda
```

```
pred_bio2 <- predict(bio_lda,X.test)  
confusionMatrix(data =pred_bio2,  
  reference = y.test)
```

```
#####PLSDA#####
```

```
set.seed(100)
```

```
plsBio <- caret::train(x = X.train,  
  y = y.train,  
  method = "pls",  
  tuneGrid = expand.grid(.ncomp = 1:15),  
  preProc = c("center","scale"),  
  metric = "Kappa",  
  trControl = ctrl)
```

```
plsBio
```

```
plot(plsBio)
```

```
predictionPLSBio <-predict(plsBio,X.test)  
confusionMatrix(data =predictionPLSBio,  
  reference =y.test)
```

```
#####Penalized Logistic Regression#####
```

```
glmGrid <- expand.grid(alpha = c( .1, .2, .4, .6, .8, 1),  
                      .lambda = seq(0, 2, length = 10))
```

```
require(caret)
```

```
set.seed(100)
```

```
bio_plr <- caret::train(X.train,  
                       y.train,  
                       method = "glmnet",  
                       metric='Kappa',  
                       tuneGrid = glmGrid,  
                       preProc = c("center", "scale"),  
                       trControl = ctrl)
```

```
bio_plr
```

```
plot(bio_plr, main='Penalized Logistic Tuning Parameters')
```

```
pred_bio4<-predict(bio_plr,X.test)
```

```
confusionMatrix(data=pred_bio4,  
                 reference=y.test)
```

```
imp<-varImp(bio_plr, scale=FALSE)
```

```
plot(imp,top=5)
```

```
### Question 12.3 ###
```

```
#install.packages("caret")
```

```
#install.packages('e1071')
```

```
library(caret)
```

```
#install.packages("modeldata")
```

```
#install.packages("generics")
```

```
library(generics)
```

```
#install.packages("tidyselect")
```

```
library(tidyselect)
```

```
library(modeldata)
```

```
data("mlc_churn")
```

```
#####NearZeroVariance#####
```

```
nZv <- nearZeroVar(mlc_churn)
```

```
length(nZv)
```

```
mlc_churn<-mlc_churn[,-nZv]
```

```
dim(mlc_churn)
```

```
###Separating predictors from the response####
```

```
churn1 <- mlc_churn[,-19]
```

```
str(churn1)
```

```
####Separating categorical from numerical variables
```

```
churn_cat <- churn1[,c(3,4,5)]
```

```
churn_num <- churn1[,~c(1,3,4,5)]
```

```
str(churn_cat)
```

```
str(churn_num)
```

```
# response... barplot
```

```
counts <- table(mlc_churn$churn)
```

```
counts
```



```
percentage <- prop.table(counts) * 100
```

```
percentage
```

```
bp <- barplot(counts,  
              names.arg = c("(Yes)", "(No)"),  
              col = c("blue", "red"),  
              main = "Churn Distribution",  
              xlab = "Churn",  
              ylab = "Count",  
              ylim = c(0, max(counts) + 10)) # Adjust y-axis limits to make space for labels
```

```
# cat... barplot
```

```
par(mfrow = c(3, 3), pin = c(2, 1))  
for (col in 1:ncol(churn_cat)) {  
  # Count the frequency of each category  
  freq <- table(churn_cat[, col])  
  barplot(freq,  
           col='lightblue',  
           main = colnames(churn_cat)[col],  
           xlab = colnames(churn_cat)[col],  
           ) # Rotates axis labels for readability  
}
```

```
# num ... histogram
```

```
par(mfrow = c(3, 3), pin = c(2, 1))  
for (col in 1:ncol(churn_num)) {  
  col_data <- churn_num[[col]] # Changed from churn_num[, col] to churn_num[[col]]
```

```

if (all(is.finite(col_data))) { # Check if all values are finite
  hist(col_data,
        col='lightblue',
        border='blue',
        main = colnames(churn_num)[col],
        xlab = colnames(churn_num)[col])
} else {
  cat("Skipping column", colnames(churn_num)[col], "due to non-finite values\n")
}
}

##outliers plot

par(mfrow = c(1, 1))

#Boxplots for all numeric variables
#numeric_data <- mlc_churn[sapply(mlc_churn, is.numeric)]
plots_per_page <- 9
for (i in seq(1, ncol(churn_num), by = plots_per_page)) {
  par(mfrow = c(3, 3))
  for (j in i:min(i+plots_per_page-1, ncol(churn_num))) {
    boxplot(churn_num[[j]],
            main = paste(names(churn_num)[j], "Distribution"),
            ylab = names(churn_num)[j])
  }
}

# correlation [cat + num]

# corr plot

```

```
correlation_matrix <- cor(churn_num)
```

```
correlation_matrix
```

```
library(corrplot)
```

```
corrplot(correlation_matrix, order = "hclust")
```

```
corrplot
```

```
#Highly correlated predictors
```

```
highCorr = findCorrelation( cor( churn_num), cutoff=0.9 )
```

```
length(highCorr)
```

```
churn_num1= churn_num[,-highCorr]
```

```
str(churn_num1)
```

```
#correlation after
```

```
cor_matrix_2 <- cor(churn_num1)
```

```
cor_matrix_2
```

```
library(corrplot)
```

```
par(mfrow = c(1,1), pin=c(2,1))
```

```
corrplot(cor_matrix_2, order = "hclust")
```

```
churn_num2 <- as.data.frame(churn_num1)
```

```
trans <- preProcess(churn_num2, method = c("BoxCox", "center", "scale", "spatialSign"))
```

```
trans
```

```
# boxplot
```

```
trfmd <- predict(trans, churn_num2)
head(trfmd)
```

```
par(mfrow = c(3, 3))
for (col in 1:ncol(trfmd)) {
  hist(trfmd[,col],
       main = colnames(trfmd)[col],
       col = 5)
}
```

```
dummRes <- dummyVars("~area_code + international_plan + voice_mail_plan",
                     data = churn_cat,
                     fullRank = TRUE)
dummy <- data.frame(predict(dummRes, newdata = churn_cat))
dummy
dim(dummy)
```

```
churn_c <- cbind(dummy, trfmd)
dim(churn_c)
str(churn_c)
###splitting####
```

```
set.seed(100)
trainR <- createDataPartition(mlc_churn$churn, p=0.8, list=FALSE)
X.train <- churn_c[trainR, ]
y.train <- mlc_churn$churn[trainR]
X.test <- churn_c[-trainR, ]
```

```
y.test <- mlc_churn$churn[-trainR]
```

```
ctrl<-trainControl(summaryFunction = twoClassSummary,classProbs = TRUE,  
  method='LGOCV',savePredictions = TRUE)
```

```
#####Logistic Regression#####
```

```
require(caret)
```

```
set.seed(100)
```

```
churn_lr <- caret::train(X.train,  
  y.train,  
  method = "glm",  
  metric = "Kappa",  
  trControl = ctrl)
```

```
churn_lr
```

```
pred_churn<-predict(churn_lr,X.test)
```

```
confusionMatrix(data=pred_churn,  
  reference=y.test)
```

```
# Predict probabilities on the test set
```

```
prob_predictions <- predict(churn_lr, X.test, type = "prob")[, 2]
```

```
# Compute the ROC curve
```

```
roc_obj <- roc(y.test, prob_predictions)
```

```
# Plot the ROC curve
```

```
plot(roc_obj, main = "ROC Curve", xlab = "False Positive Rate (1 - Specificity)", ylab = "True  
Positive Rate (Sensitivity)")
```

```
auc_value <- auc(roc_obj)
```

```
# Print the AUC value
```

```
print(auc_value)
```

```
#####Linear Discriminant Analysis#####
```

```
library(MASS)
```

```
set.seed(100)
```

```
churn_lda <- caret::train(x = X.train,
```

```
  y = y.train,
```

```
  method = "lda",
```

```
  metric = "Kappa",
```

```
  trControl = ctrl)
```

```
churn_lda
```

```
pred_churn2 <- predict(churn_lda,X.test)
```

```
confusionMatrix(data =pred_churn2,
```

```
  reference = y.test)
```

```
# Predict probabilities on the test set
```

```
prob_predictions2 <- predict(churn_lda, X.test, type = "prob")[, 2]
```

```
# Compute the ROC curve
```

```
roc_obj2 <- roc(y.test, prob_predictions2)
```

```

# Plot the ROC curve

plot(roc_obj2, main = "ROC Curve", xlab = "False Positive Rate (1 - Specificity)", ylab = "True
Positive Rate (Sensitivity)")

auc_value2 <- auc(roc_obj2)


# Print the AUC value

print(auc_value2)


#####PLSDA#####

library(MASS)
set.seed(100)


churn_plsda <- caret::train(x = X.train,
                           y = y.train,
                           method = "pls",
                           tuneGrid = expand.grid(
                             .ncomp = 1:14),
                           metric = "Kappa",
                           trControl = ctrl)

churn_plsda
plot(churn_plsda, main='PLSDA tuning parameter')
pred_churn3 <- predict(churn_plsda,X.test)
confusionMatrix(data =pred_churn3,
                 reference = y.test)

```





```
plot(churn_plr, main='Penalized Logistic Tuning Parameters')
pred_churn4<-predict(churn_plr,X.test)
```

```
confusionMatrix(data=pred_churn4,
                 reference=y.test)
```

```
library(pROC)
```

```
# Predict probabilities on the test set
```

```
prob_predictions4 <- predict(churn_plr, X.test, type = "prob")[, 2]
```

```
# Compute the ROC curve
```

```
roc_obj4 <- roc(y.test, prob_predictions4)
```

```
# Plot the ROC curve
```

```
plot(roc_obj4, main = "ROC Curve", xlab = "False Positive Rate (1 - Specificity)", ylab = "True Positive Rate (Sensitivity)")
```

```
auc_value4 <- auc(roc_obj4)
```

```
# Print the AUC value
```

```
print(auc_value4)
```